

2.1

a

A nonlinear transfer function gives the neural network a universal property: Given enough layers and neurons, the network can model any function within a certain accuracy. In a network with a linear transfer function we can only compute a linear function. A network with a linear transfer function of n layers will always be equivalent to a network with only one layer: linear functions can always be concatenated.

Whenever the function that we are trying to model isn't a linear function, it is useful to use a nonlinear transfer function. Examples include image classification or speech recognition.

b

Consider a simple neural network with two input neurons that can both either be 0 or 1 and one output layer. We want to construct an AND gate with our network, so without bias our quest would be to find $w = (w_1, w_2)$ such that:

$$\begin{aligned}0 &\leq 0 \\w_1 &\leq 0 \\w_2 &\leq 0 \\w_1 + w_2 &> 0\end{aligned}$$

which is impossible. We can easily however create the network with a bias, if we have the weights $w = (1, 1)$ and the bias $\theta = \frac{3}{2}$. Then $\text{sgn}(w^T x - \theta)$ would give us AND.

c

Point and edge filters are for example a connectionist neuron which gets values of a scalar field as input, that represent the color of each pixel or the color gradient or even a higher derivative and has weights in the following form: In the simple case of two colors (0 and 1) this point filter would return

-1	-1	-1
-1	+8	-1
-1	-1	-1

point filter

-1	0	+1
-2	0	+2
-1	0	+1

edge filter (Sobel filter)

+1	+2	+1
0	0	0
-1	-2	-1

zero for no point, 1 or -1 for a point in the outer region and 8 or -8 if the point is in the middle. This goes analogously for the other filter.

d

The first is deterministic and the second has a noise parameter and can return different states for set parameters and a given input.

2.2

The code for this assignment can be found in the file 2_2.py

a

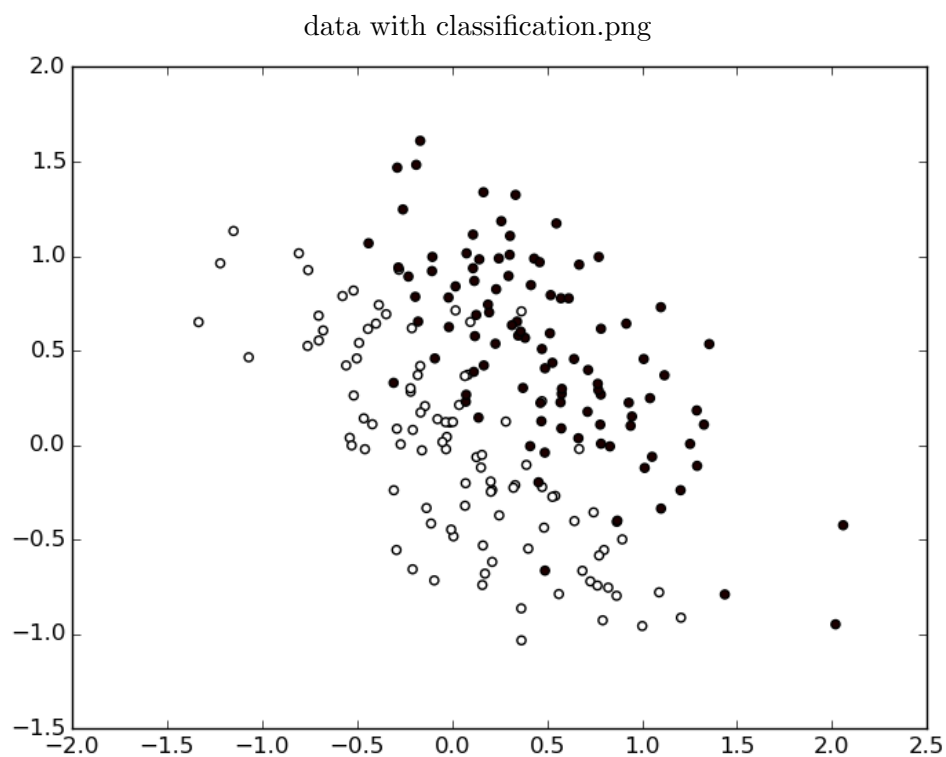


Figure 1: Input data, $x[1]$ against $x[0]$ colored according to y value

b

performance for various w-angles.png




Figure 2: Classification performance given by $\text{sign}(w \cdot x)$ for varying angles determining w

c

The weight vector giving us the best performance is $w = (0.93969262, 0.34202014)$ if we optimize it without respect to θ .

d

performance with w and theta optimized separately.png




Figure 3: Classification performance with w and θ optimized separately. The weight vector w is colored pink in the figure, and it is perpendicular to the line along which a lot of x are classified. $w * x$ gives us a measure of how far away our points, x are from this line. If we optimize w and θ one after another we get a $p = 0.805$ classification rate and an optimal $\theta = 0.135135135135$.

e

If we optimize w and θ one after another we get a $p = 0.915$ classification rate with optimal parameters $w = (0.64278761, 0.76604444)$ and $\theta = 0.339339339339$

2.3

a

A MLP could decide between a horizontal and a vertical edge, whereas a perceptionist neuron would either be able to differ between vertical edge or no vertical edge OR horizontal edge or no horizontal edge.

b




Figure 4: functions computed with normally distributed a_i with a standard deviation of 2

c




Figure 5: functions computed with normally distributed a_i with a standard deviation of 2

c + bonus




Figure 6: functions computed with normally distributed a_i with a standard deviation of 0.5




Figure 7: functions with least mean square error from $f(x) = -x$, computed with an std of 2 (green) and 0.5 (blue)